

Memory map and limits

LAB 7

Introduction and Lab objective

The objective of this lab is to introduce the tools and API for memory use, including resource limits and changing a process memory mapping.

Process and computer memory

Remember Lab 1 and the study of Linux's process memory map using the `pmap` command. You can use the shell process (usually `bash` command) as the process to study. Use `ps` to get its PID and then `pmap`¹ (you can use `pmap -xp . . .` to get the full pathname to the mapped files).

```
fso@linuxfso:~$ ps
  PID TTY          TIME CMD
  1179 pts/1        00:00:00 bash
  1923 pts/1        00:00:00 ps
fso@linuxfso:~$ pmap -x 1179
1179:  bash
Address            Kbytes      RSS   Dirty Mode  Mapping
000056064b7f8000      188      188        0 r---- bash
000056064b827000      772      768        0 r-x-- bash
000056064b8e8000      224      152        0 r---- bash
000056064b920000        16        16       16 r---- bash
000056064b924000        36        36       36 rw--- bash
000056064b92d000        44        28       28 rw--- [ anon ]
000056064d823000     2236     2156     2156 rw--- [ anon ]
00007ffbc6ed1000        64        60        0 r---- libm.so.6
00007ffbc6ee1000      460      248        0 r-x-- libm.so.6
00007ffbc6f54000      360         0        0 r---- libm.so.6
00007ffbc6fae000         4         4        4 r---- libm.so.6
00007ffbc6faf000         4         4        4 rw--- libm.so.6
00007ffbc6fb0000        28        28        0 r---- libnss_systemd.so.2
00007ffbc6fb7000      208      208        0 r-x-- libnss_systemd.so.2
[. . .]
00007ffbc7571000        40        36        0 r---- ld-linux-x86-64.so.2
00007ffbc757b000         8         8        8 r---- ld-linux-x86-64.so.2
00007ffbc757d000         8         8        8 rw--- ld-linux-x86-64.so.2
00007ffc7f9f7000      132        56       56 rw--- [ stack ]
00007ffc7fae2000        16         0        0 r---- [ anon ]
00007ffc7fae6000         8         4        0 r-x-- [ anon ]
-----
total kB              10300      6700      2428
```

Now you can better understand the mapped libraries in the process memory space, and the dynamic linker `ld-linux-x86-64.so` in the process address space. Some of those areas can possibly be shared with other processes. Can you identify the candidates? Also, some pages, are not in physical memory (look at Kbytes vs RSS). So, from approximately 42 GB of address space, this process uses aprox. 10MB virtual memory, but only aprox. 6MB are in physical memory and, from these, several MB can be shared with other processes. Should be no surprise that, if you add all the memory used by all the processes,

¹ On MacOS use `vmmap` command.

you reach a memory usage much above your computer real memory. You can check the real memory and extra swap space with the following command²:

```
fso@linuxfso:~$ free -h
```

	total	used	free	shared	buff/cache	available
Mem:	1.4Gi	576Mi	635Mi	1.9Mi	391Mi	886Mi
Swap:	974Mi	0B	974Mi			

Gi - o i potencia de 2 em vez de potencia de 10

Resource limits

Besides the OS Virtual Memory limit, each process can have stricter restrictions when executing, for memory and for other resources. Those are defined at user login for its first process and then, will be inherited by its descendants. In a terminal, check the limits for your processes by running the command:

```
fso@linuxfso:~$ ulimit -a
```

real-time non-blocking time	(microseconds, -R)	unlimited
core file size	(blocks, -c)	0
data seg size	(kbytes, -d)	unlimited
scheduling priority	(-e)	0
file size	(blocks, -f)	unlimited
pending signals	(-i)	5749
max locked memory	(kbytes, -l)	187288
max memory size	(kbytes, -m)	unlimited
open files	(-n)	1024
pipe size	(512 bytes, -p)	8
POSIX message queues	(bytes, -q)	819200
real-time priority	(-r)	0
stack size	(kbytes, -s)	8192
cpu time	(seconds, -t)	unlimited
max user processes	(-u)	5749
virtual memory	(kbytes, -v)	unlimited
file locks	(-x)	unlimited

Those are the soft limits. Some can be increased until its hard limit. Add -H to the previous command to see those. In this example, there are very few limits. Examples: 8MB for stack and 1024 open files; also, you can have a maximum of 5749 processes/threads. Run the provided memtest.c to query and print the resource limits as seen by a process:

```
fso@linuxfso:~$ memtest
```

address space limit:	soft limit=fffffffffffffffff	hard limit=fffffffffffffffff
data limit:	soft limit=fffffffffffffffff	hard limit=fffffffffffffffff
stack limit:	soft limit=	800000, hard limit=fffffffffffffffff

Change the program to create an array of 9Mbytes in the doit function (it will be allocated in the stack) and add an infinite loop at the end. Try to execute. Explain why you can't call your function in your program.

Try changing the following limits in your terminal, and run your program again:

```
fso@linuxfso:~$ ulimit -t 3
```

```
fso@linuxfso:~$ ulimit -d 10000
```

```
fso@linuxfso:~$ ulimit -s 10000
```

```
fso@linuxfso:~$ memtest
```

address space limit:	soft limit=fffffffffffffffff	hard limit=fffffffffffffffff
data limit:	soft limit=	9c4000, hard limit=9c4000
stack limit:	soft limit=	9c4000, hard limit=9c4000

Killed

Explain why it runs, the new values printed and why the infinite loop is terminated after some time.

² On MacOS you can use GUI Monitor application or `sysctl -a|grep -e swapusage -e memsize`

Testing mmap

Read the manual page and bibliography about mmap. Try the program `testing_mmap.c`. That uses mmap to map a file to its virtual memory and reads that file by reading that memory.

Sharing memory with mmap

Check the program `sharemem.c` where an anonymous memory area, created with mmap, is shared by two processes (father and son). In this example both processes read and write an array shared by both. Confirm that the father process reads values written by its child.

Readers and Writers processes

In the provided *stocks* subdirectory, there is a demo of the readers/writer pattern. Imagine a stockbroker where one program receives all the stocks transactions that several other users/programs must see. This program is in `stocks.c`, and the data to be shared is simulated by reading each line of the *stock.trace* file at some time intervals. All this data is published on a “board”, for all the other programs to read (the *clients*), using a shared memory array. Each position of the array includes a sequence number, company code, number of shares traded and its price. The implemented `client.c` program just reads the content of that array and prints to standard output. A simple check is made to the consistency of that the information. That is, using the sequence numbers, the *client* verifies if the board includes all the last 10 transactions, without missing or repeated entries. These errors can occur if a client reads the array in the middle of the update by the *stocks* program.

You need to add the needed shared memory and concurrency control implementations. Remember that you must use ‘named shared memory’ and ‘named semaphores’, as there is no father/child relations, processes can be executed at any time and there can be multiple clients running at the same time. Also, since you don’t need a file for *mmap* your shared memory, use `shm_open`. Note that, in this case, you need to use `ftruncate` after `shm_open`, to extend the shared memory to the pretended size.

Bibliography

[1] Sections about Concurrency of the recommended book, “Operating Systems: Three Easy Pieces Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau”

<http://pages.cs.wisc.edu/~remzi/OSTEP/threads-cv.pdf>

<http://pages.cs.wisc.edu/~remzi/OSTEP/threads-sema.pdf>

[2] Slides from theoretical classes (in CLIP) particularly: Aula 10 and Aula 13